MantaPay Protocol Specification
v0.4.0

Shumo Chu and Brandon H. Gomes

November 10, 2021

Abstract

MantaPay is an implementation of a *decentralized anonymous payment* scheme based on the $\text{Manta}_{\text{DAP}}$ protocol outlined in the original Manta whitepaper.

Contents

1 Introduction

TODO: add introductory remarks

2 Notation

The following notation is used throughout this specification:

- Type is the type of types[1].
- If $x : T$ then x is a value and T is a type, denoted $T : \mathsf{Type}$, and we say that x *has type* T.
- Bool is the type of booleans with values True and False.

- For any types A : Type and B : Type we denote the *type of functions* from A to B as $A \to B$: Type.
- For any types A : Type and B : Type we denote the *product type* over A and B as $A \times B$: Type with constructor $(-,-) : A \to (B \to A \times B)$.
- For any type T : Type, we define $\mathsf{Option}(T)$: Type as the inductive type with constructors:

$$\begin{aligned} \mathsf{None} &: \mathsf{Option}(T) \\ \mathsf{Some} &: T \to \mathsf{Option}(T) \end{aligned}$$

- We denote the *type of finite sets* over a type T : Type as $\mathsf{FinSet}(T)$: Type. The membership predicate for a value $x : T$ in a finite set $S : \mathsf{FinSet}(T)$ is denoted $x \in S$.
- We denote the *type of finite ordered sets* over a type T : Type as $\mathsf{List}(T)$: Type. This can either be defined by an inductive type or as a $\mathsf{FinSet}(T)$ with a fixed ordering. We denote the constructor for a list as $[\ldots]$ for an arbitrary set of elements.
- We denote the *type of distributions* over a type T : Type as $\mathfrak{D}(T)$: Type. A value x sampled from $\mathfrak{D}(T)$ is denoted $x \sim \mathfrak{D}(T)$ and the fact that the value x belongs to the range of $\mathfrak{D}(T)$ is denoted $x \in \mathfrak{D}(T)$. So namely, $y \in \{x \mid x \sim \mathfrak{D}(T)\} \leftrightarrow y \in \mathfrak{D}(T)$.
- Depending on the context, the notation $|\cdot|$ denotes either the absolute value of a quantity, the length of a list, the number of characters in a string, or the cardinality of a set.

3 Concepts

3.1 Assets

The Asset is the fundamental currency object in the MantaPay protocol. An asset a : Asset is a tuple

$$a = (a.\mathsf{id}, a.\mathsf{value}) : \mathsf{AssetId} \times \mathsf{AssetValue}$$

where the AssetId encodes the type of currency stored in a and the AssetValue encodes how many units of that currency are stored in a. MantaPay is a *decentralized anonymous payment* protocol which facilitiates the private ownership and private transfer of Asset objects.

Whenever an Asset is being used in a public setting, we simply refer to it as an Asset, but when the AssetId and/or AssetValue of a particular Asset is meant to be hidden from public view, we refer to the Asset as either, *secret*, *private*, *hidden*, or *shielded*.

Assets form the basic units of *transactions* which consume Assets on input, transform them, and return Assets on output. To preserve the economic value stored in Assets, the sum of the input AssetValues must balance the sum of the output AssetValues, and all assets in a single transaction must have the same AssetId[2]. This is called a *balanced transfer*: no AssetValue is created or destroyed in the process. The MantaPay protocol uses a distributed algorithm called Transfer to perform balanced transfers and ensure that they are valid.

3.2 Addresses

In order for MantaPay participants to send and receive Assets via the Transfer protocol, they create *addresses* which represent their partipation in the protocol. MantaPay has a 3-address system consisting of a *spending key* sk, a *viewing key* vk, and a *receiving key* rk. The keys have the following uses/properties:

- Access to a receiving key rk represents the ability to send Assets to the owner of the associated sk.
- Access to a viewing key vk represents the ability to reveal shielded Asset information for Assets belonging to the owner of the associated sk.
- Access to a spending key sk represents the ability to spend Assets that were received under the associated receiving key rk.

Participants in MantaPay are represented by their addresses, but they are not unique representations, since one participant may have access to more than one triple of keys. See § 4.2 for more information on how these keys are constructed and used for spending, viewing, and receiving Assets.

[1]By *type of types*, we mean the type of *first-level* types in some family of type universes. Discussion of the type theory necessary to make these notions rigorous is beyond the scope of this paper.

[2]It is beyond the scope of this paper to discuss transactions with inputs and outputs that feature different AssetIds, like those that would be featured in a *decentralized anonymous exchange*.

3.3 Ledger

Preserving the economic value of Assets requires more than just balanced transfers. It also requires that Assets are owned by exactly one address at a time, namely, that the ability to spend an Asset can be proved before a transfer and revoked after a transfer. It is not simply the *information-content* of an Asset that should be transfered, but the *ability to spend the asset in the future*, which should be transfered. To enforce this second invariant we can use a public ledger[3] that keeps track of the movement of Assets from one participant to another. Unfortunately, using a public ledger alone does not allow participants to remain anonymous, so MantaPay extends the public ledger by adding a special account called the *shielded asset pool* which is responsible for keeping track of the Assets which have been anonymized by the protocol. We denote the three ledger types in the protocol as follows: the public ledger as PublicLedger, the shielded asset pool as ShieldedAssetPool, and the combined ledger we denote Ledger.

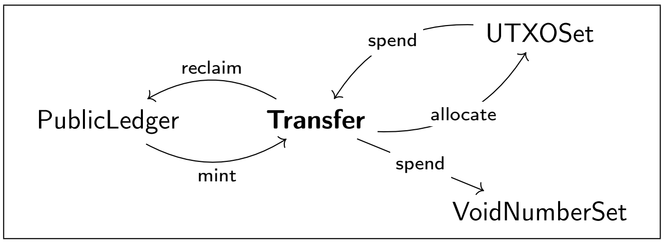

Figure 1: Lifecycle of an Asset.

The ShieldedAssetPool is made up of four parts which serve to enforce the balanced transfer of Assets among anonymous participants:

1. ShieldedAssetPool Balance: The Ledger contains a collection of Assets which encode the combined economic value of the ShieldedAssetPool and the PublicLedger. The ShieldedAssetPool balance is the subset of this total collection that has been anonymized by the MantaPay protocol. This balance is stored as a finite set of non-zero Assets.
2. § 3.3.1 UTXOSet: The UTXOSet is a collection of ownership claims to subsets of the ShieldedAssetPool (called UTXOs), each one refering to an allocated Asset transfered to a participant of the protocol.
3. § 3.3.2 EncryptedNotes: For every UTXO there is a matching EncryptedNote which contains information necessary to spend the Asset, which can be used to *provably reconstruct* the UTXO convincing the Ledger of unique ownership. The EncryptedNote can only be decrypted by the recipient of the Asset, specifically, the correct viewing key vk. See § 3.2 for more.
4. § 3.3.3 VoidNumberSet: The VoidNumberSet is a collection of commitments, like UTXOs, but which track the *spent state* of an Asset and are used to prove to the Ledger that an Asset is spent *exactly one time*.

The operation of these different parts of the ShieldedAssetPool is elaborated in the following subsections.

3.3.1 UTXOs and the UTXOSet

An *unspent transaction output*, or UTXO for short, represents a claim to the output of a balanced transfer which has otherwise *not yet been spent*. Every balanced transfer produces *public outputs*, just publicly visible Assets, and *private outputs*, represented by UTXOs, and these UTXOs are stored in the UTXOSet of the ShieldedAssetPool. A UTXO can only be claimed by the participant who owns the underlying Asset, where ownership means *knowledge of the correct spending key* and the Transfer protocol requires that all inputs to a balanced transfer *prove* that they own a UTXO which the ShieldedAssetPool has already seen in the past. The UTXOSet is *append-only* since it represents the past state of *unspent* Assets. UTXOs can only be added to the UTXOSet as outputs in the execution of a Transfer which the Ledger checks for correctness.

3.3.2 EncryptedNotes

In order to find out what Asset a UTXO is connected to, every UTXO comes with an associated EncryptedNote which stores two pieces of information, the underlying Asset, and a key diversifier, a value which allows the new owner of the Asset to reconstruct the UTXO. Being able to *provably reconstruct* a correct UTXO is a prerequisite to ownership and the ability to spend the Asset in the future. Once a participant spends an Asset that they can decrypt, they build a new EncryptedNote for the next participant that they sent their Assets to, so that they can then spend it, and so on. This is called the *in-band secret distribution*.

[3]A public (or private) ledger is not enough to solve the *provable-ownership problem* or the *double-spending problem*. A *consensus mechanism* is also required to ensure that all participants agree on the current state of the ledger. The design and specification of the consensus mechanism that secures the MantaPay ledger is beyond the scope of this paper.

3.3.3 VoidNumbers and the VoidNumberSet

Once the ability to spend an Asset is extracted from a (UTXO, EncryptedNote) pair, the ShieldedAssetPool requires another commitment in order to spend the Asset, transfering it to another participant. This commitment, called the VoidNumber, represents the revocation of the right to spend the Asset in the future, and ensures that the same Asset cannot be spent twice. Like the UTXOSet, the VoidNumberSet is *append-only* since it represents the past state of *spent* Assets. VoidNumbers can only be added to the VoidNumberSet as inputs in the execution of a Transfer which the Ledger checks for correctness.

4 Abstract Protocol

4.1 Abstract Cryptographic Schemes

In the following section, we outline the formal specifications for all of the *cryptographic schemes* used in the MantaPay protocol.

Definition 4.1.1 (Commitment Scheme)**.** A *commitment scheme* COM is defined by the schema:

$$\begin{aligned}\mathsf{Trapdoor} &: \mathsf{Type}\\ \mathsf{Input} &: \mathsf{Type}\\ \mathsf{Output} &: \mathsf{Type}\\ \mathsf{TrapdoorDistribution} &: \mathfrak{D}(\mathsf{Trapdoor})\\ \mathsf{commit} &: \mathsf{Trapdoor} \times \mathsf{Input} \to \mathsf{Output}\end{aligned}$$

with the following properties:

- **Binding**: It is infeasible to find an $x, y : \mathsf{Input}$ and $r, s : \mathsf{Trapdoor}$ such that $x \neq y$ and $\mathsf{commit}(r, x) = \mathsf{commit}(s, y)$.
- **Hiding**: For all $x, y : \mathsf{Input}$, the distributions $\{\mathsf{commit}(r, x) \mid r \sim \mathsf{TrapdoorDistribution}\}$ and $\{\mathsf{commit}(r, y) \mid r \sim \mathsf{TrapdoorDistribution}\}$ are *computationally indistinguishable*.

Notation: For convenience, we refer to $\mathsf{COM.commit}(r, x)$ by $\mathsf{COM}_r(x)$.

Definition 4.1.2 (Symmetric-Key Encryption Scheme)**.** A *symmetric-key encryption scheme* SYM is defined by the schema:

$$\begin{aligned}\mathsf{Key} &: \mathsf{Type}\\ \mathsf{Plaintext} &: \mathsf{Type}\\ \mathsf{Ciphertext} &: \mathsf{Type}\\ \mathsf{encrypt} &: \mathsf{Key} \times \mathsf{Plaintext} \to \mathsf{Ciphertext}\\ \mathsf{decrypt} &: \mathsf{Key} \times \mathsf{Ciphertext} \to \mathsf{Option}(\mathsf{Plaintext})\end{aligned}$$

with the following properties:

- **Validity**: For all keys $k : \mathsf{Key}$ and plaintexts $p : \mathsf{Plaintext}$, we have that

$$\mathsf{decrypt}(k, \mathsf{encrypt}(k, p)) = \mathsf{Some}(p)$$

- **TODO**: hiding, one-time encryption security?

Definition 4.1.3 (Key-Agreement Scheme)**.** A *key-agreement scheme* KA is defined by the schema:

$$\begin{aligned}\mathsf{PublicKey} &: \mathsf{Type}\\ \mathsf{SecretKey} &: \mathsf{Type}\\ \mathsf{SharedSecret} &: \mathsf{Type}\\ \mathsf{derive} &: \mathsf{SecretKey} \to \mathsf{PublicKey}\\ \mathsf{agree} &: \mathsf{SecretKey} \times \mathsf{PublicKey} \to \mathsf{SharedSecret}\end{aligned}$$

with the following properties:

- **Agreement**: For all $\mathsf{sk}_1, \mathsf{sk}_2 : \mathsf{SecretKey}$, $\mathsf{agree}(\mathsf{sk}_1, \mathsf{derive}(\mathsf{sk}_2)) = \mathsf{agree}(\mathsf{sk}_2, \mathsf{derive}(\mathsf{sk}_1))$
- **TODO**: security properties

Definition 4.1.4 (Key-Diversification Scheme)**.** A *key-diversification scheme* KDIV over a key-agreement scheme KA is defined by the schema:

$$\begin{aligned}\mathsf{public} &: \mathsf{KA.SecretKey} \times \mathsf{KA.PublicKey} \times \mathsf{KA.PublicKey} \to \mathsf{KA.PublicKey}\\ \mathsf{secret} &: \mathsf{KA.PublicKey} \times \mathsf{KA.SecretKey} \times \mathsf{KA.SecretKey} \to \mathsf{KA.SecretKey}\end{aligned}$$

Notation: We refer to the first argument to a KDIV function as the *diversifier* and we write it as a subscript

$$\mathsf{public}_d(x, y) := \mathsf{public}(d, x, y) \text{ and } \mathsf{secret}_d(x, y) := \mathsf{secret}(d, x, y)$$

For convenience, we also write KDIV_d to mean $\mathsf{KDIV.public}_d$ or $\mathsf{KDIV.secret}_d$, when the context is clear.

Every KDIV also has the following properties:

- **Derivation Invariance**: For any diversifier $d : \mathsf{KA.SecretKey}$ and pair of secret keys $(\mathsf{sk}_1, \mathsf{sk}_2)$ we have
$$\mathsf{KDIV}_d(\mathsf{KA.derive}(\mathsf{sk}_1), \mathsf{KA.derive}(\mathsf{sk}_2)) = \mathsf{KA.derive}(\mathsf{KDIV}_{\mathsf{KA.derive}(d)}(\mathsf{sk}_1, \mathsf{sk}_2))$$
- **TODO**: security properties?

Definition 4.1.5 (Key-Derivation Function)**.** A *key-derivation function* KDF defined over a symmetric-key encryption scheme SYM and a key-agreement scheme KA is a function of type:

$$\mathsf{KDF} : \mathsf{KA.SharedSecret} \to \mathsf{SYM.Key}$$

with the following properties:

- **TODO**: security properties

Definition 4.1.6 (Hybrid Public Key Encryption)**.** A *hybrid public key encryption scheme* [1] HPKE is an encryption scheme made up of a symmetric-key encryption scheme SYM, a key-agreement scheme KA, and a key-derivation function KDF to convert from $\mathsf{KA.SharedSecret}$ to $\mathsf{SYM.Key}$. We can define the following encryption and decryption algorithms:

- Encryption: Given a secret key $\mathsf{sk} : \mathsf{KA.SecretKey}$, a public key $\mathsf{pk} : \mathsf{KA.PublicKey}$, and plaintext $p : \mathsf{SYM.Plaintext}$, we produce the pair
$$m : \mathsf{KA.PublicKey} \times \mathsf{SYM.Ciphertext} := \big(\mathsf{KA.derive}(\mathsf{sk}), \mathsf{SYM.encrypt}(\mathsf{KDF}(\mathsf{KA.agree}(\mathsf{sk}, \mathsf{pk})), p)\big)$$
- Decryption: Given a secret key $\mathsf{sk} : \mathsf{KA.SecretKey}$, and an encrypted message, as above, $m := (\mathsf{pk}, c)$, we can decrypt m, producing the plaintext,
$$p : \mathsf{Option}(\mathsf{SYM.Plaintext}) := \mathsf{SYM.decrypt}(\mathsf{KDF}(\mathsf{KA.agree}(\mathsf{sk}, \mathsf{pk})), c)$$
which should decrypt successfully if the $\mathsf{KA.PublicKey}$ that m was encrypted with is the derived key of $\mathsf{sk} : \mathsf{KA.SecretKey}$.

Notation: We denote the above *encrypted message* type as $\mathsf{Message} := \mathsf{KA.PublicKey} \times \mathsf{SYM.Ciphertext}$, and the above two algorithms by

$$\begin{aligned}\mathsf{encrypt} &: \mathsf{KA.SecretKey} \times \mathsf{KA.PublicKey} \times \mathsf{SYM.Plaintext} \to \mathsf{Message}\\ \mathsf{decrypt} &: \mathsf{KA.SecretKey} \times \mathsf{KA.PublicKey} \times \mathsf{SYM.Ciphertext} \to \mathsf{Option}(\mathsf{SYM.Plaintext})\end{aligned}$$

TODO: security properties, combine with SYM and KA properties, like the fact that some of these keys should be ephemeral, etc.

Definition 4.1.7 (Dynamic Cryptographic Accumulator)**.** A *dynamic cryptographic accumulator* DCA is defined by the schema:

$$\begin{aligned}\mathsf{Item} &: \mathsf{Type}\\ \mathsf{State} &: \mathsf{Type}\\ \mathsf{Checkpoint} &: \mathsf{Type}\\ \mathsf{Proof} &: \mathsf{Type}\\ \mathsf{checkpoint} &: \mathsf{State} \to \mathsf{Checkpoint}\\ \mathsf{update} &: \mathsf{Item} \times \mathsf{State} \to \mathsf{State}\\ \mathsf{contains} &: \mathsf{Item} \times \mathsf{State} \to \mathsf{Option}(\mathsf{Checkpoint} \times \mathsf{Proof})\\ \mathsf{verify} &: \mathsf{Item} \times \mathsf{Checkpoint} \times \mathsf{Proof} \to \mathsf{Bool}\end{aligned}$$

with the following properties:

- **Unique Checkpoints**: For any initial state $s : \mathsf{State}$ and any list of items $I : \mathsf{List}(\mathsf{Item})$ we can generate the sequence of states:

  $$s_0 := s, \quad s_{i+1} := \mathsf{update}(I_i, s_i)$$

  Then, if we collect the checkpoints for these states, $c_i := \mathsf{checkpoint}(s_i)$, there should be exactly $|I|$-many unique checkpoints, one for each state update.
- **Provable Membership**: For any initial state $s : \mathsf{State}$ and any list of items $I : \mathsf{List}(\mathsf{Item})$ we can generate the sequences of states:

  $$s_0 := s, \quad s_{i+1} := \mathsf{update}(I_i, s_i)$$

  Then, if we collect the states s_i into a set S, we have the following property for all $s \in S$ and $t \in I$,

  $$\mathsf{Some}(c, \pi) := \mathsf{contains}(t, s), \quad \mathsf{verify}(t, c, \pi) = \mathsf{True}$$

- **TODO**: security properties

TODO: add finite capacity constraint, something like $\mathsf{update} : \mathsf{Item} \times \mathsf{State} \to \mathsf{Option}(\mathsf{State})$ where it fails when capacity is reached

Definition 4.1.8 (Non-Interactive Zero-Knowledge Proving System)**.** A *non-interactive zero-knowledge proving system* NIZK is defined by the schema:

$$\begin{aligned}
\mathsf{Statement} &: \mathsf{Type} \\
\mathsf{ProvingKey} &: \mathsf{Type} \\
\mathsf{VerifyingKey} &: \mathsf{Type} \\
\mathsf{PublicInput} &: \mathsf{Type} \\
\mathsf{SecretInput} &: \mathsf{Type} \\
\mathsf{Proof} &: \mathsf{Type} \\
\mathsf{keys} &: \mathsf{Statement} \to \mathfrak{D}(\mathsf{ProvingKey} \times \mathsf{VerifyingKey}) \\
\mathsf{prove} &: \mathsf{Statement} \times \mathsf{ProvingKey} \times \mathsf{PublicInput} \times \mathsf{SecretInput} \to \mathfrak{D}(\mathsf{Option}(\mathsf{Proof})) \\
\mathsf{verify} &: \mathsf{VerifyingKey} \times \mathsf{PublicInput} \times \mathsf{Proof} \to \mathsf{Bool}
\end{aligned}$$

Notation: We use the following notation for a NIZK:

- We write the Statement and ProvingKey arguments of prove in the superscript and subscript respectively,

  $$\mathsf{prove}^P_{\mathsf{pk}}(x, w) := \mathsf{prove}(P, \mathsf{pk}, x, w)$$

- We write the VerifyingKey argument of verify in the subscript,

  $$\mathsf{verify}_{\mathsf{vk}}(x, \pi) := \mathsf{verify}(\mathsf{vk}, x, \pi)$$

- We say that $(x, w) : \mathsf{PublicInput} \times \mathsf{SecretInput}$ has the property of being a satisfying input whenever

  $$\mathsf{satisfying}^P_{\mathsf{pk}}(x, w) := \exists \pi : \mathsf{Proof}, \mathsf{Some}(\pi) \in \mathsf{prove}^P_{\mathsf{pk}}(x, w)$$

Every NIZK has the following properties for a fixed statement $P : \mathsf{Statement}$ and keys $(\mathsf{pk}, \mathsf{vk}) \sim \mathsf{keys}(P)$:

- **Completeness**: For all $(x, w) : \mathsf{PublicInput} \times \mathsf{SecretInput}$, if $\mathsf{satisfying}^P_{\mathsf{pk}}(x, w) = \mathsf{True}$ with proof witness π, then $\mathsf{verify}_{\mathsf{vk}}(x, \pi) = \mathsf{True}$.
- **Knowledge Soundness**: For any polynomial-size adversary $\mathcal{A}$,

  $$\mathcal{A} : \mathsf{ProvingKey} \times \mathsf{VerifyingKey} \to \mathfrak{D}(\mathsf{PublicInput} \times \mathsf{Proof})$$

  there exists a polynomial-size extractor $\mathcal{E}_{\mathcal{A}}$

  $$\mathcal{E}_{\mathcal{A}} : \mathsf{ProvingKey} \times \mathsf{VerifyingKey} \to \mathfrak{D}(\mathsf{SecretInput})$$

  such that the following probability is negligible:

  $$\Pr\left[\begin{array}{c} \mathsf{satisfying}^P_{\mathsf{pk}}(x, w) = \mathsf{False} \\ \mathsf{verify}_{\mathsf{vk}}(x, w) = \mathsf{True} \end{array} \middle| \begin{array}{l} (\mathsf{pk}, \mathsf{vk}) \sim \mathsf{keys}(P) \\ (x, \pi) \sim \mathcal{A}(\mathsf{pk}, \mathsf{vk}) \\ w \sim \mathcal{E}_{\mathcal{A}}(\mathsf{pk}, \mathsf{vk}) \end{array}\right]$$

- **Statistical Zero-Knowledge**: There exists a stateful simulator $\mathcal{S}$, such that for all stateful distinguishers $\mathcal{D}$, the difference between the following two probabilities is negligible:

$$\Pr\left[\begin{array}{c}\mathsf{satisfying}^P_{\mathsf{pk}}(x,w)=\mathsf{True}\\ \mathcal{D}(\pi)=\mathsf{True}\end{array}\;\middle|\;\begin{array}{l}(\mathsf{pk},\mathsf{vk})\sim\mathsf{keys}(P)\\ (x,w)\sim\mathcal{D}(\mathsf{pk},\mathsf{vk})\\ \mathsf{Some}(\pi)\sim\mathsf{prove}^P_{\mathsf{pk}}(x,w)\end{array}\right] \text{ and } \Pr\left[\begin{array}{c}\mathsf{satisfying}^P_{\mathsf{pk}}(x,w)=\mathsf{True}\\ \mathcal{D}(\pi)=\mathsf{True}\end{array}\;\middle|\;\begin{array}{l}(\mathsf{pk},\mathsf{vk})\sim\mathcal{S}(P)\\ (x,w)\sim\mathcal{D}(\mathsf{pk},\mathsf{vk})\\ \pi\sim\mathcal{S}(x)\end{array}\right]$$

- **Succinctness**: For all $(x, w) : \mathsf{PublicInput} \times \mathsf{SecretInput}$, if $\mathsf{Some}(\pi) \sim \mathsf{prove}(P, \mathsf{pk}, x, w)$, then $|\pi| = \mathcal{O}(1)$, and $\mathsf{verify}(\mathsf{vk}, x, \pi)$ runs in time $\mathcal{O}(|x|)$.

4.2 Addresses and Key Components

Given a choice of HPKE we have the following definitions:

Definition 4.2.1 (Spending Key)**.** A SpendingKey is the following pair of keys:

$$\begin{aligned}\mathsf{view} &: \mathsf{HPKE.KA.SecretKey}\\ \mathsf{spend} &: \mathsf{HPKE.KA.SecretKey}\end{aligned}$$

The first secret key, view, is called the ViewingKey.

Definition 4.2.2 (Receiving Key)**.** A ReceivingKey is the following pair of keys:

$$\begin{aligned}\mathsf{view} &: \mathsf{HPKE.KA.PublicKey}\\ \mathsf{spend} &: \mathsf{HPKE.KA.PublicKey}\end{aligned}$$

which is derived from a spending key $\mathsf{sk} : \mathsf{SpendingKey}$ by deriving each component:

$$\begin{aligned}\mathsf{rk.view} &:= \mathsf{KA.derive}(\mathsf{sk.view})\\ \mathsf{rk.spend} &:= \mathsf{KA.derive}(\mathsf{sk.spend})\end{aligned}$$

A keypair $(\mathsf{sk}, \mathsf{rk}) : \mathsf{SpendingKey} \times \mathsf{ReceivingKey}$, represents the ability to spend and receive Assets as a unique *representative participant* on the Ledger. Any user of the MantaPay protocol can create many such keypairs, but each one represents a different participant and Assets must be transfered between them using the Transfer protocol as if they were independently owned by different users. A ReceivingKey can be used to receive any number of Assets and the SpendingKey can be used to spend any number of those Assets. See § 4.4 for the protocol used to spend a subset of Assets owned by a single user.

Important: To every spending key $\mathsf{sk} : \mathsf{SpendingKey}$ we have an assoicated viewing key $\mathsf{vk} : \mathsf{ViewingKey} := \mathsf{sk.view}$ which allows the owner to decrypt the encrypted messages associated to sk, but does not contain enough information to perform a spend with those Assets. This can be used for account auditing purposes, and for removing anonymity, but sharing this key should be done with caution.

4.3 Transfer Protocol

The Transfer protocol is the fundamental abstraction in MantaPay and facilitiates the valid transfer of Assets among participants while preserving their anonymity. The Transfer is made up of special cryptographic constructions called Senders and Receivers which represent the private input and the private output of a transaction. To perform a Transfer, a protocol participant gathers the SpendingKeys they own, selects a subset of the UTXOs they have still not spent (with a fixed AssetId), collects ReceivingKeys from other participants for the outputs, assigning each key a subset of the input Assets, and then builds a Transfer object representing the transfer they want to build. From this Transfer object, they construct a TransferPost which they then send to the Ledger to be validated and stored, representing a completed state transition in the Ledger. The transformation from Transfer to TransferPost involves keeping the parts of the Transfer that *must* be known to the Ledger and for the parts that *must* not be known, substituting them for a *zero-knowledge proof* representing the validity of the secret information known to the participant, and the Transfer as a whole.

We begin by defining the cryptographic primitives involved in the Transfer protocol:

Definition 4.3.1 (Transfer Configuration)**.** A TransferConfiguration is a collection of implementations of the following abstract cryptographic primitives:

- **Hybrid Public Key Encryption**: HPKE
- **Commitment Scheme**: COM

- **Dynamic Cryptographic Accumulator**: DCA
- **Zero-Knowledge Proving System**: NIZK

with the following notational conventions:

$$
\begin{aligned}
\mathsf{KA} &:= \mathsf{HPKE.KA} \\
\mathsf{KDIV} &:= \mathsf{HPKE.KDIV} \\
\mathsf{UTXO} &:= \mathsf{COM.Output} \\
\mathsf{VoidNumber} &:= \mathsf{COM.Output} \\
\mathsf{EncryptedNote} &:= \mathsf{HPKE.Message} \\
\mathsf{UTXOSet} &:= \mathsf{DCA}
\end{aligned}
$$

and the following constraints:

$$
\begin{aligned}
\mathsf{COM.Trapdoor} &= \mathsf{KA.PublicKey} \\
\mathsf{ValidTransfer} &: \mathsf{NIZK.Statement}
\end{aligned}
$$

where ValidTransfer is defined below.

TODO: Add the fact that COM.Input has a concatenation property. In general, add (de)serialization to the spec.
TODO: Add the fact that we have a conversion from COM.Output to KA.SecretKey, notably for diversifiers, i.e. serialize the output and then deserialize into a secret key (need to match the size)

For the rest of this section, we assume the existence of a TransferConfiguration and use the primitives outlines above explicitly. We continue by defining the Sender and Receiver constructions as well as their public counterparts, the SenderPost and ReceiverPost.

Definition 4.3.2 (Transfer Sender)**.** A Sender is the following tuple:

$$
\begin{aligned}
\mathsf{sk} &: \mathsf{SpendingKey} \\
\tilde{d} &: \mathsf{KA.PublicKey} \\
\mathsf{trapdoor} &: \mathsf{KA.PublicKey} \\
\mathsf{asset} &: \mathsf{Asset} \\
\mathsf{cm} &: \mathsf{UTXO} \\
\mathsf{cm}_c &: \mathsf{UTXOSet.Checkpoint} \\
\mathsf{cm}_\pi &: \mathsf{UTXOSet.Proof} \\
\mathsf{vn} &: \mathsf{VoidNumber}
\end{aligned}
$$

A Sender, S, is constructed from a spending key $\mathsf{sk} : \mathsf{SpendingKey}$ and an encrypted message $\mathsf{note} : \mathsf{EncryptedNote}$ with the following algorithm:

$$
\begin{aligned}
S.\mathsf{sk} &:= \mathsf{sk} \\
\tilde{d}, c &:= \mathsf{note} \\
\mathsf{Some}(\mathsf{asset}) &:= \mathsf{HPKE.decrypt}(S.\mathsf{sk.view}, \tilde{d}, c) \\
S.\mathsf{asset} &:= \mathsf{asset} \\
S.\tilde{d} &:= \tilde{d} \\
S.\mathsf{trapdoor} &:= \mathsf{KA.derive}(\mathsf{KDIV}_{S.\tilde{d}}(S.\mathsf{sk.view}, S.\mathsf{sk.spend})) \\
S.\mathsf{cm} &:= \mathsf{COM}_{S.\mathsf{trapdoor}}(S.\mathsf{asset}) \\
\mathsf{Some}(\mathsf{cm}_c, \mathsf{cm}_\pi) &:= \mathsf{UTXOSet.contains}(S.\mathsf{cm}, \mathsf{Ledger.utxos}()) \\
S.\mathsf{cm}_c &:= \mathsf{cm}_c \\
S.\mathsf{cm}_\pi &:= \mathsf{cm}_\pi \\
S.\mathsf{vn} &:= \mathsf{COM}_{S.\tilde{d}}(S.\mathsf{sk.view} \,||\, S.\mathsf{sk.spend})
\end{aligned}
$$

Definition 4.3.3 (Transfer Sender Post)**.** A SenderPost is the following tuple extracted from a Sender:

$$
\begin{aligned}
\mathsf{cm}_c &: \mathsf{UTXOSet.Checkpoint} \\
\mathsf{vn} &: \mathsf{VoidNumber}
\end{aligned}
$$

which are the parts of a Sender which should be *posted* to the Ledger.

Definition 4.3.4 (Transfer Receiver)**.** A $\mathsf{Receiver}$ is the following tuple:

$$\begin{aligned} \mathsf{rk} &: \mathsf{ReceivingKey} \\ d &: \mathsf{KA.SecretKey} \\ \mathsf{trapdoor} &: \mathsf{KA.PublicKey} \\ \mathsf{asset} &: \mathsf{Asset} \\ \mathsf{cm} &: \mathsf{UTXO} \\ \mathsf{note} &: \mathsf{EncryptedNote} \end{aligned}$$

A $\mathsf{Receiver}$, R, is constructed from a receving key $\mathsf{rk} : \mathsf{ReceivingKey}$, an asset $\mathsf{asset} : \mathsf{Asset}$, and a given[4] diversifier $d : \mathsf{HPKE.KA.SecretKey}$ with the following algorithm:

$$\begin{aligned} R.\mathsf{rk} &:= \mathsf{rk} \\ R.d &:= d \\ R.\mathsf{trapdoor} &:= \mathsf{KDIV}_{R.d}(R.\mathsf{rk.view}, R.\mathsf{rk.spend}) \\ R.\mathsf{asset} &:= \mathsf{asset} \\ R.\mathsf{cm} &:= \mathsf{COM}_{R.\mathsf{trapdoor}}(R.\mathsf{asset}) \\ R.\mathsf{note} &:= \mathsf{HPKE.encrypt}(R.d, R.\mathsf{rk.view}, \mathsf{asset}) \end{aligned}$$

Definition 4.3.5 (Transfer Receiver Post)**.** A $\mathsf{ReceiverPost}$ is the following tuple extracted from a $\mathsf{Receiver}$:

$$\begin{aligned} \mathsf{cm} &: \mathsf{UTXO} \\ \mathsf{note} &: \mathsf{EncryptedNote} \end{aligned}$$

which are the parts of a $\mathsf{Receiver}$ which should be *posted* to the Ledger.

Definition 4.3.6 (Transfer Sources and Sinks)**.** A Source (or a Sink) is an Asset representing a public input (or output) of a $\mathsf{Transfer}$.

Definition 4.3.7 (Transfer Object)**.** A $\mathsf{Transfer}$ is the following tuple:

$$\begin{aligned} \mathsf{sources} &: \mathsf{List}(\mathsf{Asset}) \\ \mathsf{senders} &: \mathsf{List}(\mathsf{Sender}) \\ \mathsf{receivers} &: \mathsf{List}(\mathsf{Receiver}) \\ \mathsf{sinks} &: \mathsf{List}(\mathsf{Asset}) \end{aligned}$$

The *shape* of a $\mathsf{Transfer}$ is the following 4-tuple of cardinalities of those sets

$$\left(|T.\mathsf{sources}|, |T.\mathsf{senders}|, |T.\mathsf{receivers}|, |T.\mathsf{sinks}|\right)$$

In order for a $\mathsf{Transfer}$ to be considered *valid*, it must adhere to the following constraints:

- **Same Id**: All the $\mathsf{AssetIds}$ in the $\mathsf{Transfer}$ must be equal.
- **Balanced**: The sum of input $\mathsf{AssetValues}$ must be equal to the sum of output $\mathsf{AssetValues}$.
- **Well-formed Senders**: All of the $\mathsf{Senders}$ in the $\mathsf{Transfer}$ must be constructed according to the above Sender definition.
- **Well-formed Receivers**: All of the $\mathsf{Receivers}$ in the $\mathsf{Transfer}$ must be constructed according to the above $\mathsf{Receiver}$ definition.

In order to prove that these constraints are satisfied for a given $\mathsf{Transfer}$, we build a zero-knowledge proof which will witness that the $\mathsf{Transfer}$ is valid and should be accepted by the Ledger. It is not necessary to prove that the encryption of $\mathsf{Receiver.note}$ and the decryption of a note from the Ledger are valid. Deviation from the protocol in encryption or decryption stages does not reduce the security of the protocol for honest participants.

Definition 4.3.8 (Transfer Validity Statement)**.** A transfer $T : \mathsf{Transfer}$ is considered *valid* if and only if

[4] The key-diversifier is not directly chosen by the ledger participants building the $\mathsf{Transfer}$. Instead, it is derived from other $\mathsf{Transfer}$ data and the current state of the ledger. See Def 4.3.8 for more.

1. All the $\mathsf{AssetIds}$ in T are equal:

$$\left|\left(\bigcup_{a\in T.\mathsf{sources}} a.\mathsf{id}\right) \cup \left(\bigcup_{S\in T.\mathsf{senders}} S.\mathsf{asset.id}\right) \cup \left(\bigcup_{R\in T.\mathsf{receivers}} R.\mathsf{asset.id}\right) \cup \left(\bigcup_{a\in T.\mathsf{sinks}} a.\mathsf{id}\right)\right| = 1$$

2. The sum of input $\mathsf{AssetValues}$ is equal to the sum of output $\mathsf{AssetValues}$:

$$\left(\sum_{a\in T.\mathsf{sources}} a.\mathsf{value}\right) + \left(\sum_{S\in T.\mathsf{senders}} S.\mathsf{asset.value}\right) = \left(\sum_{R\in T.\mathsf{receivers}} R.\mathsf{asset.value}\right) + \left(\sum_{a\in T.\mathsf{sinks}} a.\mathsf{value}\right)$$

3. For all $S \in T.\mathsf{senders}$, the Sender S is well-formed:

$$\begin{aligned} S.\mathsf{trapdoor} &= \mathsf{KA.derive}(\mathsf{KDIV}_{S.\tilde{d}}(S.\mathsf{sk.view}, S.\mathsf{sk.spend})) \\ S.\mathsf{cm} &= \mathsf{COM}_{S.\mathsf{trapdoor}}(S.\mathsf{asset}) \\ S.\mathsf{vn} &= \mathsf{COM}_{S.\tilde{d}}(S.\mathsf{sk.view} \,||\, S.\mathsf{sk.spend}) \\ \mathsf{UTXOSet.verify}(S.\mathsf{cm}, S.\mathsf{cm}_c, S.\mathsf{cm}_\pi) &= \mathsf{True} \end{aligned}$$

4. For all $(i, R) \in \mathsf{enumerate}(T.\mathsf{receivers})$, the $\mathsf{Receiver}$ R is well-formed at index i with respect to FAIR:

$$\begin{aligned} R.d &= \mathsf{COM}_{\mathsf{FAIR}}(i \,||\, R.\mathsf{rk.view} \,||\, R.\mathsf{rk.spend}) \\ R.\mathsf{trapdoor} &= \mathsf{KDIV}_{R.d}(R.\mathsf{rk.view}, R.\mathsf{rk.spend}) \\ R.\mathsf{cm} &= \mathsf{COM}_{R.\mathsf{trapdoor}}(R.\mathsf{asset}) \end{aligned}$$

   where FAIR is the following constant, binding the new diversifiers to the current state of the ledger:

$$\mathsf{FAIR} := \mathsf{COM}_f(\mathsf{UTXOSet.checkpoint}(\mathsf{Ledger.utxos}()) \,||\, \mathsf{Concat}_{S\in T.\mathsf{senders}}(S.\mathsf{sk.spend}))$$

   and $f \sim \mathsf{COM.TrapdoorDistribution}$ is a randomly chosen trapdoor.

Notation: This statement is denoted $\mathsf{ValidTransfer}$ and is assumed to be expressible as a $\mathsf{Statement}$ of NIZK.

Definition 4.3.9 (Transfer Post)**.** A $\mathsf{TransferPost}$ is the following tuple:

$$\begin{aligned} \mathsf{sources} &: \mathsf{List}(\mathsf{Source}) \\ \mathsf{senders} &: \mathsf{List}(\mathsf{SenderPost}) \\ \mathsf{receivers} &: \mathsf{List}(\mathsf{ReceiverPost}) \\ \mathsf{sinks} &: \mathsf{List}(\mathsf{Sink}) \\ \pi &: \mathsf{NIZK.Proof} \end{aligned}$$

A $\mathsf{TransferPost}$, P, is constructed by assembling the zero-knowledge proof of $\mathsf{Transfer}$ validity from a known proving key $\mathsf{pk} : \mathsf{NIZK.ProvingKey}$ and a given $T : \mathsf{Transfer}$:

$$\begin{aligned} x &:= \mathsf{Transfer.public}(T) \\ w &:= \mathsf{Transfer.secret}(T) \\ \mathsf{Some}(\pi) &\sim \mathsf{NIZK.prove}_{\mathsf{pk}}^{\mathsf{ValidTransfer}}(x, w) \\ P.\mathsf{sources} &:= x.\mathsf{sources} \\ P.\mathsf{senders} &:= x.\mathsf{senders} \\ P.\mathsf{receivers} &:= x.\mathsf{receivers} \\ P.\mathsf{sinks} &:= x.\mathsf{sinks} \\ P.\pi &:= \pi \end{aligned}$$

where $\mathsf{Transfer.public}$ returns $\mathsf{SenderPosts}$ for each Sender in T and $\mathsf{ReceiverPosts}$ for each $\mathsf{Receiver}$ in T, keeping $\mathsf{Sources}$ and Sinks as they are, and $\mathsf{Transfer.secret}$ returns all the rest of T which is not part of the output of $\mathsf{Transfer.public}$.

Now that a participant has constructed a transfer post $P : \mathsf{TransferPost}$ they can send it to the Ledger for verification.

Definition 4.3.10 (Ledger-side Transfer Validity)**.** To check that P represents a valid Transfer, the ledger checks the following:

- **Public Withdraw**: All the public addresses corresponding to the Assets in P.sources have enough public balance (i.e. in the PublicLedger) to withdraw the given Asset.
- **Public Deposit**: All the public addresses corresponding to the Assets in P.sinks exist.
- **Shielded Withdraw**: The total balance in P.sinks does not exceed the amount in the ShieldedAssetPool balance.
- **Current Checkpoint**: The UTXOSet.Checkpoint stored in each P.senders is equal to current checkpoint, UTXOSet.checkpoint(Ledger.utxos()), for the current state of the Ledger.
- **New VoidNumbers**: All the VoidNumbers in P.senders are unique, and no VoidNumber in P.senders has already been stored in the Ledger.VoidNumberSet.
- **New UTXOs**: All the UTXOs in P.receivers are unique, and no UTXO in P.receivers has already been stored on the ledger.
- **Verify Transfer**: Check that $\mathsf{NIZK.verify}_{\mathsf{vk}}(P.\mathsf{sources} \,||\, P.\mathsf{senders} \,||\, P.\mathsf{receivers} \,||\, P.\mathsf{sinks}, P.\pi) = \mathsf{True}$.

Definition 4.3.11 (Ledger Transfer Update)**.** After checking that a given TransferPost P is valid, the Ledger updates its state by performing the following changes:

- **Public Updates**: All the relevant public accounts on the PublicLedger are updated to reflect their new balances using the Sources and Sinks present in P.
- **Pool Update**: The ShieldedAssetPool balance is updated to reflect the new shielded balances, increasing by the amount:
$$\left(\sum_{a \in P.\mathsf{sources}} a.\mathsf{value}\right) - \left(\sum_{a \in P.\mathsf{sinks}} a.\mathsf{value}\right)$$
- **UTXOSet Update**: The new UTXOs are appended to the UTXOSet.
- **VoidNumberSet Update**: The new VoidNumbers are appended to the VoidNumberSet.

4.4 Semantic Transactions

For MantaPay participants to use the Transfer protocol, they will need to keep track of the current state of their shielded assets and use them to build TransferPosts to send to the Ledger. The *shielded balance* of any participant is the sum of the balances of their shielded assets, but this balance may be fragmented into arbitrarily many pieces, as each piece represents an independent asset that the participant received as the output of some Transfer. To then spend a subset of their shielded balance, the participant would need to accumulate all of the relevant fragments into a large enough *shielded asset* to spend all at once, building a collection of TransferPosts to send to the Ledger.

Any wallet implementation should see that their users need not keep track of this complexity themselves. Instead, like a public ledger, the notion of a *transaction* between one participant and another should be viewed as a single action that the user can take, performing a withdrawl from their shielded balance. To describe such a *semantic transaction*, we assume the existence of two transfer shapes[5]: Mint with shape $(1, 0, 1, 0)$ and PrivateTransfer with shape $(0, N, N, 0)$ for some natural number $N > 1$.

For a fixed spending key, $\mathsf{sk} : \mathsf{SpendingKey}$, and asset id, $\mathsf{id} : \mathsf{AssetId}$, we are given a balance state, $\mathcal{B} : \mathsf{FinSet}\,(\mathsf{KA.PublicKey} \times \mathsf{AssetValue})$, a set of diversifier-balance pairs for unspent assets, a total balance to withdraw, $\mathsf{total} : \mathsf{AssetValue}$, and a receiving key $\mathsf{rk} : \mathsf{ReceivingKey}$. We can then compute

$$\textsc{BuildTransaction}(\mathsf{sk}, \mathcal{B}, \mathsf{total}, \mathsf{rk})$$

to receive a List(TransferPost) to send to the ledger, representing the transfer of total to rk.

If all of the Transfers are accepted by the ledger, the balance state $\mathcal{B}$ should be updated accordingly, removing all of the pairs which were used in the Transfer.

[5]Other Transfer accumulation algorithms are possible with different starting shapes.

Algorithm 1 Semantic Transaction Algorithm

```
procedure BuildTransaction(sk, 𝓑, total, rk)
    B ← Sample(total, 𝓑)                                  ▷ Samples pairs from 𝓑 that total at least total
    if len(B) = 0 then
        return [ ]                                        ▷ Insufficient Balance
    end if
    P ← [ ]                                               ▷ Allocate a new list for TransferPosts
    while len(B) > N do                                   ▷ While there are enough pairs to make another Transfer
        A ← [ ]
        for b ∈ (B, N) do                                 ▷ Get the next N pairs from B
            S ← BuildSenders_sk(b)
            [acc, zs...] ← BuildAccumulatorAndZeroes_sk(S)  ▷ Build a new accumulator and zeroes
            P ← P + TransferPost(Transfer([], S, [acc, zs...], []))
            (A, Z) ← (A + (acc.d̃, acc.asset.value), Z + zs) ▷ Save acc for the next loop, zs for the end
        end for
        B ← A + remainder(B, N)
    end while
    S ← PrepareZeroes_sk(N, B, Z, P)                      ▷ Use Z and Mints to make B go up to N in size.
    R ← BuildReceiver_sk(rk, S)
    [c, zs...] ← BuildAccumulatorAndZeroes_sk(S)
    return P + TransferPost(Transfer([], S, [R, c, zs...], []))
end procedure
```

5 Concrete Protocol

TODO: other than cryptographic schemes, are there any implementation details we want to include here?

5.1 Concrete Cryptographic Schemes

Definition 5.1.1 (Commitment Scheme)**.** For COM, we use the *Pedersen Commitment Scheme* with 256 windows of size 4 each.

Definition 5.1.2 (Symmetric-Key Encryption Scheme)**.** For SYM, we use

TODO: symmetric-key encryption scheme: AES-GCM with magic-number nonce and no associated data

Definition 5.1.3 (Key-Agreement Scheme)**.** For KA, we use

TODO: key-agreement scheme: x25519 elliptic curve Diffie-Hellman key exchange

Definition 5.1.4 (Key-Diversification Scheme)**.** For KDIV, we use

TODO: key-diversification scheme: extension of key-agreement

Definition 5.1.5 (Key-Derivation Function)**.** For KDF, we use

TODO: key-derivation function: Blake2s with magic-number salt

Definition 5.1.6 (Hybrid Public Key Encryption)**.** For HPKE, we use

TODO: integrated encryption scheme: combination of above encryption protocols

Definition 5.1.7 (Dynamic Cryptographic Accumulator)**.** For DCA, we use

TODO: dynamic cryptographic accumulator: Merkle Tree with Pedersen hashes (incremental tree for the ledger is an optimization since it only needs to know enough to compute the checkpoint as values are accumulated)

Definition 5.1.8 (Non-Interactive Zero-Knowledge Proofs)**.** For NIZK, we use

TODO: non-interactive zero-knowledge proving system: Groth16 and/or PLONK

6 Differences from $\text{MANTA}_{\text{DAP}}$

6.1 Reusable Addresses

TODO: compare old one-time address protocol to reusable addresses and why reusable is better

6.2 Transfer Circuit Unification

TODO: compare new single transfer circuit to the many old circuits

7 Acknowledgements

TODO: add acknowledgements

8 References

References

[1] Richard Barnes, Karthikeyan Bhargavan, Benjamin Lipp, and Christopher A. Wood. Hybrid Public Key Encryption. Internet-Draft draft-irtf-cfrg-hpke-12, Internet Engineering Task Force, September 2021. Work in Progress.