

MantaPay Trusted Setup Protocol Specification

v0.0.0

Francisco Hernández Iglesias, Todd Norton *

September 26, 2022

Abstract

We describe the protocol for the MantaPay trusted setup ceremony to generate prover and verifier keys for Groth16 ZK-SNARK proofs.

Contents

1	Introduction	2
2	Context	2
2.1	Circuit	2
2.2	Quadratic Arithmetic Programs	2
2.3	The Groth16 Setup function	3
2.4	Multi-Party Computation	3
2.5	Phase Structure	4
2.5.1	Phase 1	4
2.5.2	Phase 2	4
3	Requirements	4
3.1	Goals	4
3.2	Non-Goals	5
4	Design	5
4.1	Ceremony Protocol	5
4.2	Messaging Protocol	6
4.3	Server State Machine	6
4.4	Client State Machine	7
5	References	8

*ordered alphabetically.

1 Introduction

The MantaPay protocol (ref. to the specs) guarantees transaction privacy by using the Groth16 [2] Non-Interactive Zero-Knowledge Proving System (NIZK). In short, Groth16 is defined over a bilinear pairing of elliptic curves $e : \mathbb{G}_1 \times \mathbb{G}_2 \rightarrow \mathbb{G}_\tau$ over a prime field \mathbb{F}_p . Let $\phi \in \mathbb{F}_p^\ell$ denote the set of *public inputs*, $w \in \mathbb{F}_p^{m-\ell}$ the set of *witnesses*, whose knowledge we want to prove, and let $\tau \in (\mathbb{F}_p^*)^4$ be a set of randomly generated numbers known as the *simulation trapdoor*. Groth16 consists of four parts:

- $(\sigma, \tau) \leftarrow \text{Setup}$: Randomly generates τ , from which it computes σ , which consists of elliptic curve points in \mathbb{G}_1 and \mathbb{G}_2 . σ is to be understood as the proving and verifying keys for Groth16.
- $\pi \leftarrow \text{Prove}(\sigma, \phi, w)$: Computes a proof of knowledge of w , π , for a given setup σ and public input ϕ .
- $0, 1 \leftarrow \text{Verify}(\sigma, \phi, \pi)$: Checks whether the proof π is valid against the setup σ and the public input ϕ .
- $\pi \leftarrow \text{Sim}(\tau, \phi)$: Simulates a proof that will always be valid when verified against the setup σ corresponding to τ and the public input ϕ .

It is important to note that the Sim function is what makes the Groth16 protocol zero-knowledge: you can compute a valid proof π for any given setup σ and public input ϕ without knowledge of the witness w , provided that you have access to the simulation trapdoor τ . But Sim also makes Groth16 potentially insecure: if a malicious agent knew τ for a given σ , they could fabricate valid proofs for any statement regardless of its veracity.

The goal of the *trusted setup* is to provide a Groth16 setup σ in a secure way, i.e., in such a way that nobody has access to the trapdoor τ that was used to compute it.

2 Context

2.1 Circuit

Throughout this paper, by circuit we mean a *Rank-1 Constraint System (R1CS)*. It is defined as a system of equations over \mathbb{F}_r of the form

$$\sum_{i=0}^m a_i u_{i,q} \cdot \sum_{i=0}^m a_i v_{i,q} = \sum_{i=0}^m a_i w_{i,q}, \quad q = 1, \dots, n, \quad (1)$$

where $a_0 = 1$. This system of equations, in the context of zero-knowledge proofs, is to be understood as follows:

- The numbers $u_{i,q}, v_{i,q}, w_{i,q}$ are constants in \mathbb{F}_r which represent the operations performed in an arithmetic circuit. Here constant means constant in the protocol, e.g. MantaPay will have a fixed set of $u_{i,q}, v_{i,q}, w_{i,q}$.
- The numbers $\phi = (a_1, \dots, a_\ell)$ are the public inputs. In MantaPay, these correspond to the TransferPost, excluding the proof.
- The numbers $w = (a_{\ell+1}, \dots, a_m)$ are the witnesses. In MantaPay, these correspond to the elements of the Transfer which are not part of the TransferPost.
- An R1CS defines the following binary relation

$$R = \left\{ (\phi, w) \mid \phi = (a_1, \dots, a_\ell), w = (a_{\ell+1}, \dots, a_m), (1) \text{ is satisfied} \right\} \subset \mathbb{F}_r^\ell \times \mathbb{F}_r^{m-\ell} \quad (2)$$

- The statements that can be proved in this terminology are of the form: for a given circuit (1) and public input ϕ , there exists¹ a witness w such that $(\phi, w) \in R$.

2.2 Quadratic Arithmetic Programs

Quadratic Arithmetic Programs (QAPs) give an alternative way to describe a circuit, equivalent to R1CS. A QAP is a system of polynomial equations of the form

$$\sum_{i=0}^m a_i u_i(X) \cdot \sum_{i=0}^m a_i v_i(X) \equiv \sum_{i=0}^m a_i w_i(X) \pmod{t(X)}, \quad (3)$$

where

¹and I know

- $u_i(X), v_i(X), w_i(X) \in \mathbb{F}_r[X]$ are degree $n - 1$ polynomials, and $t(X) \in \mathbb{F}_r[X]$ is a degree n polynomial, all of which are fixed for the protocol.
- The numbers $\phi = (a_1, \dots, a_\ell)$ are the public inputs.
- The numbers $w = (a_{\ell+1}, \dots, a_m)$ are the witnesses.
- A QAP defines the following binary relation

$$R = \left\{ (\phi, w) \mid \phi = (a_1, \dots, a_\ell), w = (a_{\ell+1}, \dots, a_m), (3) \text{ is satisfied} \right\} \subset \mathbb{F}_r^\ell \times \mathbb{F}_r^{m-\ell} \quad (4)$$

- The statements that can be proved in this terminology are of the form: for a given circuit (3) and public input ϕ , there exists a witness w such that $(\phi, w) \in R$.

We derive the QAP description of a circuit from its R1CS description as follows:

1. Choose k as the minimal integer such that $2^k \geq n$. Fix $t(X) = X^{2^k} - 1$.
2. We derive the polynomial $u_i(X)$ from the R1CS vector $(u_{i,q})_{q=1}^n$ via a Lagrange basis $\{L_q(x)\}_{q \in Q}$ for the domain I . That is,

$$u_i(X) = \sum_{q=1}^n u_{i,q} L_q(x) \quad (5)$$

3. The reader may readily check that (1) is equivalent to (3) with these definitions.

2.3 The Groth16 Setup function

Let us now recall how the Groth16 Setup function works in the MantaPay protocol. We start with

- The pairing curve BN254 [?], which consists of a triple of elliptic curves $(\mathbb{G}_1, \mathbb{G}_2, \mathbb{G}_T)$ and a non-degenerate bilinear map $e : \mathbb{G}_1 \times \mathbb{G}_2 \rightarrow \mathbb{G}_T$. We fix generators g and h of \mathbb{G}_1 and \mathbb{G}_2 , respectively. Note that $\text{ord}(g) = \text{ord}(h) = r$ is prime, and that $e(g, h)$ generates a subgroup of order r of $\mathbb{G}_T \cong F_{p^{12}}^*$.
- The MantaPay circuit, encoded as a QAP $\{u_i(X), v_i(X), w_i(X), t(X)\}$.

All public parameters derive from a simulation trapdoor $\tau = (\alpha, \beta, \delta, x) \leftarrow \mathbb{F}_r^*$ (the famous “toxic waste”). The Groth16 public parameters themselves are elements of \mathbb{G}_1 and \mathbb{G}_2 , specifically

$$\sigma_1 = \left[\left(\alpha, \beta, \delta, \{x^i\}_{i=0}^{n-1}, \{\beta u_i(x) + \alpha v_i(x) + w_i(x)\}_{i=0}^\ell, \left\{ \frac{\beta u_i(x) + \alpha v_i(x) + w_i(x)}{\delta} \right\}_{i=\ell+1}^m, \left\{ \frac{x^i t(x)}{\delta} \right\}_{i=0}^{n-2} \right) \right]_1 \quad (6)$$

$$\sigma_2 = \left[\left(\beta, \delta, \{x^i\}_{i=0}^{n-1} \right) \right]_2 \quad (7)$$

where $[y]_1 = y \cdot g$ and $[y]_2 = y \cdot h$ for all $y \in \mathbb{F}_r$.

The output of Setup is $\sigma = (\sigma_1, \sigma_2)$. These are the public parameters from which Groth16 proofs are formed. The trapdoor τ is *not* public; indeed, a malicious prover with knowledge of τ could construct fraudulent proofs. The goal of the trusted setup is to generate σ without revealing τ .

2.4 Multi-Party Computation

A decentralized way to generate σ without revealing τ is to compute σ in such a way that τ becomes a shared secret split among a diverse set of participants. This may be achieved via *secure multi-party computation* (MPC). The MPC we employ is a protocol for computing σ incrementally from private inputs τ_i belonging to participants P_i in the computation.

The key security property of this MPC is that its security is ensured by having at least one honest participant. An honest participant is one who keeps their private input τ_i from all other participants, ideally by permanently clearing it from their system’s memory after participation. Put differently, this *1-out-of-N* honest participants guarantee states that to determine the toxic waste τ requires the collusion of *all* participants in the MPC.

By soliciting contributions to the MPC from a diverse set of participants, we increase the difficulty of such collusion. Note that any individual with a stake in the security of MantaPay can guarantee this personally, simply by participating honestly in the Setup MPC.

(? Maybe mention verifiability of the MPC transcript?)

2.5 Phase Structure

The full Setup MPC splits usefully into two *phases*. In *Phase 1* we generate a modified KZG setup (todo cite KZG) which is *universal* in the sense that these parameters may be used by any ZK circuit of small enough size. In *Phase 2* we derive σ from the output of Phase 1. The parameters generated in Phase 2 are *circuit-specific*: they depend on the QAP description of the circuit (3) and must be computed separately for each ZK circuit. This two-phase splitting of the MPC is formalized in [1].

2.5.1 Phase 1

The first class consists of those which only depend on the elliptic curves and not on the circuit, i.e., on the QAP. These parameters, namely (TODO: Get the number of powers of $[x]_1$ right)

$$KZG = (\{[x]_1\}_{i=0}^{2n-2}, \{[\alpha x]_1\}_{i=0}^{n-1}, \{[\beta x]_1\}_{i=0}^{n-1}, [\beta]_2, \{[x]_2\}_{i=0}^{n-1}) \quad (8)$$

are known as the *Kate-Zaverucha-Goldberg (KZG)* commitments. Since these parameters don't depend on the specifics of the circuit, we take them from the Perpetual Powers of Tau (PPoT) project [?], a multi-party computation similar to our trusted setup described below, with a 1-out-of- N security assumption. To ensure its soundness, we have personally verified each contribution to PPoT.

2.5.2 Phase 2

The rest of the Groth16 parameters do depend on the circuit, so we have the corresponding values for the MantaPay circuit. The rest of the paper is devoted to explaining in detail the multi-party computation known as the trusted setup ceremony which we perform to make sure the remaining Groth16 parameters are computed safely.

- Refer to above definition of σ
- Explain initialization: at least mention that inputs are (8) plus QAP coefficients.
- Mention that at this point we have prover keys but with $\delta = 1$, and the point is to modify δ .
- Outline how σ is derived from these inputs. This part is MPC, reference protocol description below.

Also need to define what is a phase 2 Contribution, what is a Proof. (Maybe this is in Protocol Design?)

3 Requirements

3.1 Goals

- Anonymity: Pre-registration only requires a Twitter handle and an email address, none of which needs to be linked to your identity.
- Server coordination: The ceremony will be coordinated by a server, which will:
 - Manage the contribution queue for the participants, supporting priority tiers.
 - Ensure that only pre-registered participants with valid signatures are allowed to take part in the ceremony.
 - Ensure that no more than one participant is contributing at a given time.
 - Ensure that no participant contributes more than once.
 - Inform the participants about their queue/contribution status.
 - Check the validity proof of the proposed contributions before adding them to the ceremony.
 - After a successful contribution, return its hash to the participant.
- Updatable: More contributions can be added to the ceremony at later stages if desired.
- Verifiable: The ceremony and all its contributions can be verified by independent auditors. We ensure that is the case by:
 - Publishing the transcript of the ceremony, including the validity proof and hash of each contribution.

- Asking participants to publish their hashes in independent locations, e.g. Twitter.
- Distributing an open-source verification library.

3.2 Non-Goals

- **Permissionless:** We require participants to pre-register to contribute to the ceremony. However, registration is open to anyone with a Twitter account and an email address.
- **Support for independently computed contributions:** The ceremony only supports those contributions done through the server with our client. However, the server code is open-source and the transcript of the ceremony is publicly available for audits.

4 Design

4.1 Ceremony Protocol

Here we describe the MPC protocol for the trusted setup ceremony. This is a protocol for N contributors C_1, \dots, C_N to participate in the computation of the public parameters σ . Because the participants must contribute in serial, we use a central **Coordinator** to order the participants and check their contributions. We emphasize that the **Coordinator** exists solely to help organize the ceremony and has no effect on the security properties of the MPC.

The protocol is initialized by the **Coordinator** and proceeds in repeated rounds, in which the **Coordinator** and a **Contributor** exchange messages. The messaging protocol is described in Section 4.2.

Initialization

The **Coordinator** computes an initial state σ_0 and challenge point c_0 as

$$(\sigma_0, c_0) = \text{initialize}(\text{circuits, KZG parameters})$$

The inputs to `initialize` are the R1CS descriptions (1) of some ZK circuits and modified KZG parameters (8) (the output of a prior Phase 1 ceremony). The function `initialize` computes σ_0 , a Groth16 proving key with $\delta = 1$ (see (6)), and an initial challenge point c_0 .

Note that `initialize` is deterministic and the circuit descriptions and KZG parameters are public. Therefore any observer may verify that **Coordinator** completed this step correctly.

Contribution Round

The n^{th} round of contribution begins with the state σ_{n-1} and challenge c_{n-1} of the previous round. **Contributor** C_n will interact with **Coordinator** to produce the next state σ_n and a proof π_n that it was computed according to the protocol. The round consists of these steps:

1. **Coordinator** sends (σ_{n-1}, c_{n-1}) to **Contributor** C_n .
2. **Contributor** computes $(\sigma_n, \pi_n) = \text{contribute}(\sigma_{n-1}, c_{n-1})$. The `contribute` function must:
 - (a) Sample $\delta_n \in \mathbb{F}$ randomly
 - (b) Compute contribution (TODO: ref. def.)
 - (c) Generate proof (TODO: ref. def.)
 - (d) Remove δ_n from memory
3. **Contributor** sends (σ_n, π_n) to **Coordinator**
4. **Coordinator** computes `verify` $(\sigma_{n-1}, c_{n-1}, \sigma_n, \pi_n)$. This checks that the **Contributor** has formed σ_n from σ_{n-1} according to the protocol.
 - (a) If the check fails, **Coordinator** rejects this contribution. Nothing is added to the transcript and the **Coordinator** proceeds to next round with state and challenge unchanged ($\sigma_n = \sigma_{n-1}$ and $c_n = c_{n-1}$).
 - (b) Otherwise, **Coordinator** computes challenge $c_n = \text{challenge}(\sigma_{n-1}, c_{n-1}, \sigma_n, \pi_n)$ and records σ_n, π_n, c_n to the Transcript.
5. **Coordinator** proceeds to next round with (σ_n, c_n) .

This process repeats until all Contributors have made their contribution. At the end (assuming all contributions were valid) we have a Groth16 prover key σ with $\delta = \delta_1 \cdot \delta_2 \cdot \dots \cdot \delta_n$ and a Transcript $T = \{(\sigma_i, \pi_i, c_i)\}_{i=1}^N$ recording all contributions to the ceremony and allowing a third-party to verify that σ was computed according to protocol.

4.2 Messaging Protocol

In each round, the Coordinator and Contributor communicate via (unencrypted?) messages. Messages from the Contributor to the Coordinator are signed with an Ed25519 signature. The Coordinator accepts only those messages with a valid signature whose public verifying key belongs to a Registry of participants. The Registry and signature checks prevent a DDoS attack on the ceremony in which malicious participants fill up the contribution queue and intentionally time-out without contributing.

The Contributor sends one of two messages to the Coordinator: a QueryRequest or an UpdateRequest. Each message follows the format

(Participant ID | Domain Tag | Nonce | Payload).

In a QueryRequest the Payload is empty, whereas in an UpdateRequest the Payload contains a new state and proof, (σ_n, π_n) . In both cases the message is serialized to bytes and signed with an Ed25519 signature.

The Coordinator responds to these messages according to the current state of the protocol:

- QueryRequest: the Coordinator responds with a QueryResponse containing:
 - Current state σ_{n-1} and challenge c_{n-1} , if Contributor is at front of Queue (see below).
 - Contributor’s current position in Queue, if Contributor is not at front of queue.
 - Error message, if Contributor has already participated in ceremony.
- UpdateRequest: the Coordinator responds with an UpdateResponse informing the Contributor of whether their contribution was successfully verified.
- If the message from Contributor fails signature verification, the Coordinator responds with the expected nonce for the Contributor’s Participant ID. (If the Participant ID is invalid the Coordinator ignores the message.)

4.3 Server State Machine

The Coordinator role is performed by a central server state machine. The Coordinator’s duties are to enforce the MPC protocol and organize the contributions.

To enforce the MPC protocol, the Coordinator performs:

- Parameter Initialization: a reproducible initial state σ_0 and initial challenge c_0 are computed from public data.
- Contribution Verification: each contribution to the ceremony is checked to conform to the MPC protocol.
- Contribution Archival: each successful contribution to the ceremony is recorded, together with the proof that it conforms to the MPC protocol.

To organize the contributions, the Coordinator performs:

- Registry Maintenance: a registry of participants records who is authorized to submit contributions to the ceremony and the public keys with which they will sign their messages.
- Signature Verification: all messages from a Contributor to the Coordinator will be signed.
- Queue Management: during the ceremony, participants are ordered in a queue.

State

The above Coordinator tasks are accomplished by a machine whose state consists of:

- Registry: For each participant, a record of their public signing key, current signature nonce, and whether they have already contributed to the ceremony.
- MpcState: The current pair (σ_n, c_n) (see above).
- Transcript: The history of MPC states and proofs (see above).

- **Queue:** an ordering of the participants waiting to contribute. This may be a priority queue, if desired.²
- **TimedLock:** a lock is given to the participant at the front of the queue while they compute their contribution. This lock times out after a specified duration and drops the participant from the queue.

State Changes

The following function calls change the state of the machine:

- **initialize:** set **MpcState** to (σ_0, c_0) (see above).
- **enqueue(ParticipantId):** check that participant is registered and has not already contributed. If so, add to end of **Queue**.
- **update:** Compute **verify** $(\sigma_{n-1}, c_{n-1}, \sigma_n, \pi_n)$ to check that the latest contribution conforms to MPC protocol. If so,
 - compute challenge c_n , update **MpcState** to (σ_n, c_n)
 - Add (σ_n, π_n, c_n) to **Transcript**
 - Set participant's contribution status in **Registry**.
 - Update **Queue** and **TimedLock**.

If contribution is invalid, update only **Queue** and **TimedLock** (and downgrade participant's **Queue** priority, if applicable).

Operation

The **Coordinator** state machine initializes its state and then listens for messages and processes these according to the messaging protocol (see above). Two types of messages are recognized, **QueryRequest** and **UpdateRequest**. The state machine responds in the following way to each request:

- **QueryRequest**
 1. Parse **ParticipantId** from request.
 2. Consult **Registry**, confirm participant is registered and has not contributed.
 3. Consult **Queue**; if participant is not in **Queue**, call **enqueue**.
 4. Consult **Queue**; if participant is at front of **Queue**, send a message containing **MpcState** to participant and set **TimedLock**. Otherwise, send participant a message containing their current position in **Queue**.
- **UpdateRequest**
 1. Parse **ParticipantId** from request. Check that the **TimedLock** refers to this participant.
 2. Parse state, proof (σ_n, π_n) from request.
 3. Call **update** $(\sigma_{n-1}, c_{n-1}, \sigma_n, \pi_n)$.

If any of the above checks do not pass, the **Coordinator** responds with an appropriate error message.

TODO: Flow chart would express this nicely

4.4 Client State Machine

The duty of a **Contributor** is to contribute according to the MPC protocol. To this end, a client state machine performs:

- **Ed25519 Keypair Generation:** each **Contributor** generates their own signature credentials.
- **Message Signing:** messages to **Coordinator** are signed.
- **Contribution:** formed according to MPC protocol.

²A priority queue can help defend against a DDoS attack in which malicious participants intentionally time out without contributing. If a participant times out too many times they can be de-prioritized, allowing more reliable participants to skip them in the queue.

State

The state of this machine is an immutable Ed25519 private key. The machine also needs access to some random oracle to perform keypair generation and contribution. (TODO: The state also includes a signer, since it keeps track of nonces the state can change)

This state is initialized by `generate-keypair`, which generates an Ed25519 private/public keypair. Note that the same state can later be recovered by prompting users to enter their private key or seed phrase.

Operation

1. Initialize state.
2. Send signed `QueryRequest` to `Coordinator`. Await response containing `MpcState` and challenge, (σ_{n-1}, c_{n-1}) .
3. Compute $(\sigma_n, \pi_n) = \text{contribute}(\sigma_{n-1}, c_{n-1})$.
 - (a) Sample random δ
 - (b) Transform σ_{n-1} by δ to get σ_n
 - (c) Produce ratio proof π_n
4. Send signed `UpdateRequest` to `Coordinator`. Await response confirming submission.
5. Clear δ from memory.

Steps 3a and 5 of this protocol have famously led to participants sampling δ using randomness from Chernobyl and later smashing their computers to clear it from memory. An ambitious participant can easily modify our implementation (todo: cite code) to accommodate exotic sources of randomness.

5 References

References

- [1] Sean Bowe, Ariel Gabizon, and Ian Miers. Scalable multi-party computation for zk-snark parameters in the random beacon model. Cryptology ePrint Archive, Paper 2017/1050, 2017. <https://eprint.iacr.org/2017/1050>.
- [2] Jens Groth. On the size of pairing-based non-interactive arguments. In *EUROCRYPT (2)*, volume 9666 of *Lecture Notes in Computer Science*, pages 305–326. Springer, 2016.